

Mini Project 12 — Securities Investor Grievance Routing Service

Course: UE24AM342AA5 — ML System Design and AgentOps | **Domain:** Capital Markets and Investor Protection
Component: Jackfruit (Level 3), 10 marks | **Team size:** 2–3 students

Technical terms used here are defined in 00_GLOSSARY.pdf.

1. Context

The Securities and Exchange Board of India (SEBI) operates a centralised investor complaint system. Complaints arrive as free text from retail investors about brokers, mutual funds, registrars, listed companies and depositories.

The first decision is jurisdictional. Not every complaint written to SEBI belongs to SEBI. A payment that failed at the bank end belongs to the Reserve Bank of India's Ombudsman. **Sending a complaint to the wrong regulator delays redress.**

The second decision is which regulated entity is responsible, because that fixes the redressal timeline. **A single complaint can involve two entities, and sometimes two regulators.**

2. What to Build

A multi-agent service that accepts an investor complaint, validates it, determines jurisdiction and the responsible entity from the regulations, and routes it with a cited justification.

This is a production service, not a notebook.

The system never disposes of a complaint. It produces a routing determination with supporting evidence. A grievance officer decides.

3. Corpus

SEBI publishes all circulars and regulations free on its website. The RBI Integrated Ombudsman Scheme is a free PDF. All documents below are free to download.

- SEBI's master circular on investor grievance redressal, including the matters it covers
- The SEBI (Mutual Funds) Regulations, the chapter on investor servicing
- The RBI Integrated Ombudsman Scheme

Prefer documents with a text layer. A scanned document needs optical character recognition (OCR) before it can be indexed.

Why this corpus is difficult. Each regulator's framework lists what it covers, not what it does not. A complaint that uses investment vocabulary about a bank product will be retrieved against the wrong framework.

4. What the System Must Do

#	Requirement
F1	Accept an investor complaint over a documented REST interface, against a published data contract
F2	Enforce schema, range, categorical and cross-field validation at the ingestion boundary . Quarantine rejects.
F3	Classify the complaint domain - broker, mutual fund, registrar, listed company, depository, out-of-jurisdiction
F4	Retrieve the jurisdictional provisions and the servicing obligations of the entity identified
F5	Produce a jurisdiction and responsibility determination grounded in retrieved provisions , with circular-level citations
F6	Route to one of: assign-to-entity with timeline, refer-to-other-regulator, or not-maintainable with reason. State the basis.

#	Requirement
F7	Escalate to a human instead of answering when the complaint splits across regulators, or the entity type falls outside your corpus
F8	Expose health, metrics and structured logs for the whole pipeline

F5 and F7 hold most of the difficulty. A system that always answers is easy to build; one that knows when not to answer is not.

4.1 The difficult case

The system will be tested against this scenario at the demonstration.

"I invested in a fixed deposit through my bank's app. The interest credited is lower than what I was promised. I want to complain to SEBI."

The complaint uses investment vocabulary and names SEBI, so a search returns SEBI's grievance framework. But a bank fixed deposit is a banking product, and complaints about banks fall under the RBI Integrated Ombudsman Scheme. The system must establish the product and the entity, find the framework that covers them, and route the complaint to the correct regulator with citations.

At Evaluation 1, show the passages in your corpus that this case depends on, and the values you will test with. Once confirmed, this is the reference answer for your demonstration.

5. Design Approach

The straightforward implementation embeds the request, retrieves the highest-scoring passages, and asks the model to answer from them. This handles ordinary cases but not the case in section 4.1.

The reason is structural. A single retrieval pass returns passages resembling the question asked. The case in section 4.1 turns on a provision the request never mentions, so a single embedding-based search is unlikely to return it.

The required approach is an agent that:

1. Retrieves the provision that appears to apply
2. Determines whether any other provision modifies it, and searches for such provisions
3. Searches again, using a query derived from what it has read
4. Verifies that the retrieved passages support the answer
5. Retries with a different strategy, or escalates, if they do not

Steps 2 and 4 distinguish an agent from a fixed pipeline. Both are assessed.

6. Agent System Requirements

Unit 4 — agent architecture, orchestration, containerisation, observability.

#	Requirement
A1	Build the reasoning layer with LangGraph or CrewAI . Retrieval must be a tool the agent chooses to call , not a fixed stage. Provide at least three tools, such as corpus search, lookup of a clause by identifier, and lookup of the request's own fields. Bound the loop by steps, tool calls and wall-clock time.
A2	At least two agents with distinct responsibilities and a defined hand-off. Document the topology and what happens if one fails mid-run. Two prompts in one process are not two agents.
A3	Verify every answer against the passages retrieved, before returning it. Every citation must resolve to retrieved text. If verification fails, search again with a different strategy or escalate. An unverified answer must not be returned. This is the basis of F7.
A4	Expose the service with FastAPI . Runs take tens of seconds, so execute them asynchronously : return a run identifier immediately, and the caller polls a status endpoint or receives a callback. Justify the choice. Place a queue or task runner between submission and execution.

#	Requirement
A5	Timeouts on every model and tool call, bounded retry with exponential backoff , a circuit breaker on the provider, and graceful degradation along a defined fallback path. All must be demonstrable on request. A run that did not complete must not be returned as a confident answer.
A6	Track per agent and display on the dashboard: task success rate , steps and tool calls per run, latency end to end and per step, retry rate , container uptime.
A7	<code>docker compose up --scale <agent-worker>=3</code> must work. Measure the effect on throughput and p95, using the stub described in D6. If scaling does not help, identify the actual bottleneck.
A8	A simple web client that submits a request, obtains the result, and displays the recommendation with citations and the escalation path. No marks for appearance.

7. LLMOps Requirements

Unit 3 — *retrieval, prompt management, evaluation, guardrails, monitoring.*

#	Requirement
L1	Orchestrate retrieval with LangChain or LlamaIndex ; index with FAISS or Chroma . Document the chunking and embedding strategy and how it addresses the difficulty in section 3. Re-embedding must be possible without rebuilding the whole index.
L2	Treat prompts as versioned artifacts , and record the prompt version with every response. Compare at least two prompt variants on the same evaluation set; report which performed better, on which metric, and by how much.
L3	Build a fixed evaluation set of at least 20 scenarios , each with its expected outcome and the passage that supports it, including the case in section 4.1, runnable as a batch from the command line. Use TruLens or equivalent. Report groundedness, context relevance, faithfulness and hallucination rate , stating how each is computed. Then make one deliberate change and report all four metrics before and after .
L4	Prompt injection defence on every free-text field reaching a model, output filtering so the system cannot issue a determination it is not authorised to make, and toxicity detection on input and output. Agents are more exposed than pipelines, because injected text can reach a model holding tools.
L5	Record prompt tokens, completion tokens and cost at the time of the call , attributed per step as well as per run . Report cost per 100 runs and the largest single cost driver.
L6	A feedback endpoint allowing the caller to mark a result satisfactory or not, with a reason, persisted and aggregated. Score a small sample of live runs automatically against the evaluation metrics, charting quality over time.

8. Platform and Deployment Requirements

Units 1 and 2 — *data contracts, architecture, containerisation, observability, reliability.*

#	Requirement
D1	Publish a data contract for every external interface. Enforce schema, type, range, category and cross-field validation at the ingestion boundary , before any agent starts. Quarantine rejected records with their field-level errors . A malformed request must not start an agent run.
D2	An idempotent pipeline that reads, parses, chunks, embeds and indexes the corpus, invocable from a command-line tool. Record data freshness — when the index was built, from which versions, how many chunks — and the pipeline's throughput, and expose both as metrics.
D3	At least four services under Docker Compose, starting with a single <code>docker compose up</code> from a clean clone, with dependencies and health checks declared. Nothing runs outside the containers. Justify the placement of deployment boundaries on scaling, failure isolation, independent release or resource profile.
D4	No credentials, endpoints or model identifiers anywhere in the repository . All environment-specific configuration injected at runtime, with a documented <code>.env.example</code> .

#	Requirement
D5	Structured logging throughout, a correlation ID across every service and agent step, tracing with OpenTelemetry exported to Jaeger or Grafana Tempo, Prometheus scraping every service, and a Grafana dashboard showing p50/p95 latency, error rate, throughput, availability and resource utilisation per container. Given a run identifier, the entire run must be reconstructible — every tool call, query, passage, prompt version, per-step latency and cost, and the verification outcome.
D6	Declare a p95 service level objective (SLO) for a complete run before measuring anything , justified against section 1. Measure a run from submission to final result, not the submission call alone. Load test with Locust , with the LLM replaced by a stub service that returns fixed responses after a realistic delay, so that the test measures your system rather than the provider's rate limits. Report RPS, p50 and p95 against the target and state whether it was met. Identify the concurrency at which the target fails and the component responsible . A target set after seeing the results does not satisfy this requirement.
D7	Caching — of embeddings, retrievals or whole runs, justified — and rate limiting per client. Report the measured effect of caching on p95 and cost per run. Rate limiting is a cost control here: one client looping against a multi-step agent can exhaust the budget in minutes.
D8	A runbook for the two most likely failure modes — symptom, confirmation, recovery — and an incident log with at least one genuine entry recording a real failure during development.

9. Demonstration

Prepare all ten items.

#	Item
1	docker compose up from a clean clone to a working system
2	A valid request producing a routed recommendation with citations that resolve to retrieved text
3	One run traced end to end — every tool call, query, passage, prompt version, per-step latency and cost
4	A malformed request rejected and quarantined, with evidence that no agent run started
5	A case the corpus cannot answer producing an escalation rather than a fabricated response
6	Self-verification catching an unsupported answer, followed by a further search or escalation
7	A prompt injection attempt and the guardrail response
8	The circuit breaker opening when the provider is unavailable, and the degradation path
9	Locust results against the declared target, the concurrency at which it fails, and the effect of - - scale
10	The case in section 4.1, handled correctly

10. Assessment

If the system does not start and serve a valid request at the demonstration, the project is awarded zero marks, irrespective of the report, presentation or repository.

Establish a working end-to-end path by week two and add requirements after that. Completing all 22 requirements is not expected; sound prioritisation is itself assessed.

Criterion	Marks
Problem understanding, architecture and objective clarity	10
Presentation, teamwork and communication	10
Engineering judgement — tool selection, scope, handling of constraints	10
Report, code quality and repository documentation	10
Total (Evaluation 2, scaled to 8 marks)	40

11. Deliverables

Deliverable	Due
Corpus confirmation and project plan	Evaluation 1
Report and final GitHub repository	One week before the presentation
Live demonstration and viva	Evaluation 2

Report — a plain report covering, briefly: problem and scope; agent design (topology, tools, verification, loop limits); an architecture diagram with deployment boundaries marked; design decisions and what was rejected; evaluation results before and after the change; the p95 target, measured results, bottleneck and cost per 100 runs; and what did not work.

Repository — the README must be sufficient for someone unfamiliar with the project to run the system without assistance. This will be tested. Include the runbook, incident log, `.env.example` and evaluation set.

12. Cost

All software and data required are free. The corpus documents are public; **sentence-transformers run locally for embeddings at no cost**, and a small model such as all-MiniLM-L6-v2 keeps memory use low; FAISS, Chroma, LangGraph, CrewAI, LangChain, LlamaIndex, TruLens, FastAPI, Redis, Prometheus, Grafana, OpenTelemetry, Jaeger and Locust are all free. For the model, use a provider's free tier. A local model via Ollama is also possible but needs a considerably more powerful machine.

Embed locally. Embedding a corpus through a paid API is the only operation here with a real cost.

Agents multiply model calls — typically several per run, against one for a fixed pipeline. Free tiers also have low rate limits, which is why load testing uses a stub (D6). This is the reason for the loop limits in A1, the caching in D7 and the per-step cost tracking in L5. Set a spending limit on the account before starting, and place the model behind an interface you control so the provider can be changed by configuration.

13. Out of Scope

- Model training is not required.
- Cloud deployment is not required; the system runs locally under Docker Compose.
- **Use synthetic complaints only.** Do not submit anything to any real regulator portal.